

Exercise 4: Solution

Loss: BCE - Forward method

```
def forward(self, y_out, y_truth):
    """
    Performs the forward pass of the binary cross entropy loss function.

    :param y_out: [N, ] array predicted value of your model.
        y_truth: [N, ] array ground truth value of your training set.
    :return: [N, ] array of binary cross entropy loss for each sample of your training set.
    """
    result = None

    #####
    # TODO:                                     #
    # Implement the forward pass and return the output of the BCE loss. #
    #####

    result = -y_truth * np.log(y_out) - (1 - y_truth) * np.log(1 - y_out)

    #####
    #                                     END OF YOUR CODE                                     #
    #####

    return result
```

Loss: BCE – Backward method

```
def backward(self, y_out, y_truth):  
    """  
    Performs the backward pass of the loss function.  
  
    :param y_out: [N, ] array predicted value of your model.  
        y_truth: [N, ] array ground truth value of your training set.  
    :return: [N, ] array of binary cross entropy loss gradients w.r.t y_out for  
        each sample of your training set.  
    """  
    gradient = None  
  
    #####  
    # TODO:                                     #  
    # Implement the backward pass. Return the gradient wrt y_out    #  
    #####  
  
    gradient = - (y_truth / y_out) + (1 - y_truth) / (1 - y_out)  
  
    #####  
    #                                     END OF YOUR CODE                                     #  
    #####  
    return gradient
```

Classifier: Sigmoid

```
def sigmoid(self, x):  
    """  
    Computes the output of the sigmoid function  
  
    :param x: input of the sigmoid, np.array of any shape  
    :return: output of the sigmoid with same shape as input vector x  
    """  
    out = None  
  
    #####  
    # TODO:                                     #  
    # Implement the sigmoid function, return out #  
    #####  
  
    out = 1 / (1 + np.exp(-x))  
  
    #####  
    #                                     END OF YOUR CODE #  
    #####  
  
    return out
```

Classifier: Forward method

```
def forward(self, X):
    """
    Performs the forward pass of the model.

    :param X: N x D array of training data. Each row is a D-dimensional point.
    :return: Predicted labels for the data in X, shape N x 1
             1-dimensional array of length N with classification scores.
    """
    assert self.W is not None, "weight matrix W is not initialized"
    # add a column of 1s to the data for the bias term
    batch_size, _ = X.shape
    X = np.concatenate((X, np.ones((batch_size, 1))), axis=1)
    # save the samples for the backward pass
    self.cache = X
    # output variable
    y = None

    #####
    # TODO:                                     #
    # Implement the forward pass and return the output of the model. Note   #
    # that you need to implement the function self.sigmoid() for that       #
    #####

    y = X.dot(self.W)
    y = self.sigmoid(y)

    #####
    #                                     END OF YOUR CODE                 #
    #####

    return y
```

Classifier: Backward method

```
def backward(self, dout):
    """
    Performs the backward pass of the model.

    :param dout: N x M array. Upstream derivative. It is as the same shape of the forward() output.
    | | | If the output of forward() is z, then it is dL/dz, where L is the loss function.
    :return: dW --> Gradient of the weight matrix, w.r.t the upstream gradient 'dout'.
    """
    assert self.cache is not None, "Run a forward pass before the backward pass. Also, don't forget to st
    | such as in 'self.cache = (X, y, ...)"
    dW = None

    #####
    # TODO:
    # Implement the backward pass. Return the gradient w.r.t W --> dW.
    # Make sure you've stored ALL needed variables in self.cache.
    #
    # Hint 1: It is recommended to follow the Stanford article on
    # calculating the chain-rule, while dealing with matrix notations:
    # http://cs231n.stanford.edu/handouts/linear-backprop.pdf
    #
    # Hint 2: Remember that the derivative of sigmoid(x) is independent of x,
    # and could be calculated with the result, calculated earlier at
    # the forward() function.
    #####

    # We calculate the derivatives in order, like in the chain rule.
    # Let us denote y = XW + b, z = sigmoid(y)

    X, y = self.cache

    # 1) dL/dy = dL/dz * dz / dy. According to stanford's trick:
    dz_dy = y * (1 - y)
    dl_dy = dout * dz_dy # Now, this is the upstream derivative for step 2.

    # 2) dL/dw = dL/dy * dy/dw. According to stanford's trick:
    dW = X.T.dot(dl_dy)

    #####
    # END OF YOUR CODE
    #####
    return dW
```

Keep the dimensions of the arrays
in mind:

X: [N, D + 1]

W: [D+1, M]

y: [N, M] (Sigmoid over XW)

dW should be of shape [D + 1, M] as
it contains a gradient of the output
w.r.t. W for each sample. The
average over all samples is taken in
the solver step.

Optimization

Optimizer: Step method

```
def step(self, dw):  
    """  
    :param dw: [D+1,1] array gradient of loss w.r.t weights of your linear model  
    :return weight: [D+1,1] updated weight after one step of gradient descent  
    """  
  
    weight = self.model.W  
  
    #####  
    # TODO:                                                    #  
    # Implement the gradient descent for 1 step to compute the weight  #  
    #####  
  
    weight -= self.lr * dw  
  
    #####  
    #                               END OF YOUR CODE                               #  
    #####  
  
    self.model.W = weight
```


Solver: Step method

```
def _step(self):
    """
    Make a single gradient update. This is called by train() and should not
    be called manually.
    """
    model = self.model
    loss_func = self.loss_func
    X_train = self.X_train
    y_train = self.y_train
    opt = self.opt
    #####
    #   TODO:                                     #
    #   Get the gradients dhat{y}/dW and dLoss/dhat{y}.           #
    #   Combine them via the chain rule to obtain dLoss / dW.     #
    #   Proceed by performing an optimizing step using the given  #
    #   optimizer (by calling opt.step() with the gradient wrt W). #
    #                                                             #
    #   Hint 1: Don't forget the order of operations: forward, loss, #
    #   backward.                                                  #
    #####
    model_forward = model.forward(X_train)
    _, loss_grad = loss_func(model_forward, y_train)
    grad = model.backward(loss_grad)
    opt.step(grad)
    #####
    #                               END OF YOUR CODE                #
    #####
```

Model and loss_func return (forward, backward) when called, cf. `__call__()` in their base classes.

Mind the dimensions of all elements. In particular, we want to update W (via `opt.step()`) with an array of the same shape, i.e., $[D + 1, M]$

Questions? Piazza