Exercise 11: Sentiment analysis

Implement Embedding

```python
class Embedding(nn.Module):
    def __init__(self, num_embeddings, embedding_dim, padding_idx):
        """
        """
        super().__init__()
        self.num_embeddings = num_embeddings
        self.embedding_dim = embedding_dim

        # We handle the padding for you
        self.padding_idx = padding_idx
        self.register_buffer(
            'padding_mask',
            (torch.arange(0, num_embeddings) != padding_idx).view(-1, 1)
        )

        self.weight = None
        ##################################################################
        # TODO: Set self.weight to a parameter intialized with standard normal #
        # N(0, 1) and has a shape of (num_embeddings, embedding_dim).     #
        ##################################################################
        self.weight = nn.Parameter(torch.randn(num_embeddings, embedding_dim))
        ##################################################################
        #                      END OF YOUR CODE                          #
        ##################################################################

        # Handle the padding
        self.weight.data[padding_idx] = 0
```

```python
    def forward(self, inputs):
        """
        Inputs:
            inputs: A long tensor of indices of size (seq_len, batch_size)
        Outputs:
            embeddings: A float tensor of size (seq_len, batch_size, embedding_dim
        """

        # Ensure <eos> always return zeros
        # and padding gradient is always 0
        weight = self.weight * self.padding_mask

        embeddings = None

        ##################################################################
        # TODO: Select the indices in the inputs from the weight tensor       #
        # hint: It is very short                                         #
        ##################################################################
        embeddings = weight[inputs]
        ##################################################################
        #                      END OF YOUR CODE                          #
        ##################################################################

        return embeddings
```

Here we build Embedding
for generating representations.

Implement an RNN Classifier

```python
class RNNClassifier(nn.Module):
    def __init__(self, num_embeddings, embedding_dim, hidden_size, use_lstm=True, **additional_kwargs):
        """
        """
        super().__init__()
        # Change this if you edit arguments
        self.hparams = {
            'num_embeddings': num_embeddings,
            'embedding_dim': embedding_dim,
            'hidden_size': hidden_size,
            'use_lstm': use_lstm,
            **additional_kwargs
        }
        ########################################################################
        # TODO:                                                                #
        ########################################################################
        self.embedding = Embedding(num_embeddings, embedding_dim, 0)
        self.rnn = (LSTM if use_lstm else RNN)(embedding_dim, hidden_size)
        self.output = nn.Linear(hidden_size, 1)
        ########################################################################
        #                          END OF YOUR CODE                           #
        ########################################################################
```

- Here Embedding will be passed to the recurrent neural network.

Implement an RNN Classifier

```python
class RNNClassifier(nn.Module):
    def __init__(self, num_embeddings, embedding_dim, hidden_size, use_lstm=True, **additional_kwargs):

    def forward(self, sequence, lenghts=None):
        """
        """

        output = None


        ########################################################################
        # TODO: Apply the forward pass of your network                         #
        # hint: Don't forget to use pack_padded_sequence if lenghts is not None #
        # packed_padded_sequence should be applied to the embedding outputs    #
        ########################################################################
        embeddings = self.embedding(sequence)
        if lenghts is not None:
            embeddings = pack_padded_sequence(embeddings, lenghts)

        h_seq, h = self.rnn(embeddings)
        if isinstance(h, tuple):
            h = h[0]

        output = self.output(h.squeeze(0)).sigmoid().view(-1)
        ########################################################################
        #                         END OF YOUR CODE                             #
        ########################################################################

        return output
```

- Sequences need to have the same length
- For this, we use the function *pack_padded_sequence* to truncate longer reviews and pad shorter reviews with zeros .

Implement an RNN Classifier

```python
class RNNClassifier(nn.Module):
    def __init__(self, num_embeddings, embedding_dim, hidden_size, use_lstm=True, **additional_kwargs):

    def forward(self, sequence, lenghts=None):
        """
        """
        output = None

        ####################################################################
        # TODO: Apply the forward pass of your network                     #
        # hint: Don't forget to use pack_padded_sequence if lenghts is not None   #
        # packed_padded_sequence should be applied to the embedding outputs     #
        ####################################################################
        embeddings = self.embedding(sequence)
        if lenghts is not None:
            embeddings = pack_padded_sequence(embeddings, lenghts)

        h_seq, h = self.rnn(embeddings)
        if isinstance(h, tuple):
            h = h[0]

        output = self.output(h.squeeze(0)).sigmoid().view(-1)
        ####################################################################
        #                        END OF YOUR CODE                         #
        ####################################################################

        return output
```

- After the RNN layer, we use sigmoid as the activation function as our task is a binary prediction.

Create model

```python
from exercise_code.rnn.tests import classifier_test, parameter_test
from exercise_code.rnn.text_classifiers import RNNClassifier
model = None
####################################################################
# TODO - Create a Model                                            #
####################################################################
model = RNNClassifier(len(vocab), 22, 32)
####################################################################
#                       END OF YOUR CODE                           #
####################################################################

# Check whether your model is sufficiently small and have a correct output format
parameter_test(model), classifier_test(model, len(vocab))
```

```
Total number of parameters: 117245
Your model is sufficiently small :)
All output tests are passed :)!

(True, True)
```

Remark:
Try to use other *embedding dimensions* and *hidden sizes* for better results.

Training

```python
###################################################################
# TODO - Train Your Model                                         #
###################################################################
optim = torch.optim.Adam(model.parameters())
epochs = 5
gclip = 40
# Training loop
for e in range(epochs):
    print('Epoch {}...\n'.format(e))
    print('Starting training...')
    model.train()
    num_corrects = 0
    num_labels = 0
    total_loss = 0.0
    for i, data in enumerate(train_loader):
        seq = data['data'].to(device)
        label = data['label'].to(device)
        seq_lens = data['lengths']

        model.zero_grad()

        pred = model(seq, seq_lens)
        loss = bce_loss(pred, label)
        loss.backward()
        clip_grad_norm_(model.parameters(), max_norm=gclip)
        optim.step()

        num_corrects += ((pred > 0.5) == label).sum().item()
        num_labels += label.numel()
        total_loss += loss.item() * label.numel()

        if i > 0 and i % 100 == 0:
            print('Step {} / {}, Loss {}'.format(i, len(train_loader), loss.item()))

    print('Training Loss: {}, Training Accuracy: {}'.format(
        total_loss / num_labels, num_corrects / num_labels
    ))
    print('\nStarting evaluation...')
    model.eval()
    print('Evaluation Accuracy:', compute_accuracy(model, val_loader))
    print('\n')
```

- As we have learned from the lecture, RNN may suffer from exploding gradients. To tackle this problem, we use a technique here called gradient clipping with the function *clip_grad_norm_* .

It rescales the gradients to keep them small and avoid taking a huge descent step.

Parameter tuning

Embedding dimension	Hidden size	gclip	epochs	Accuracy on test set
32	32	30	8	0.8426
50	40	30	8	0.8410
40	30	30	8	0.8508
43	33	30	8	0.8478
40	30	40	8	0.8430
40	30	40	5	0.8365

With this architecture, we achieved an accuracy of 85.1% with the parameters above.

Optional exercise: recurrent neural network

Implement LSTM

```python
class LSTM(nn.Module):
    def __init__(self, input_size=1, hidden_size=20):
        super().__init__()
        """
        """

        self.hidden_size = hidden_size
        self.input_size = input_size

        ########################################################################
        # TODO:
        ########################################################################
        self.W_hh = nn.Linear(self.hidden_size, 4 * self.hidden_size, bias=True)
        self.W_xh = nn.Linear(self.input_size, 4 * self.hidden_size, bias=True)
        ########################################################################
        #                          END OF YOUR CODE                          #
        ########################################################################
```

In the initialization of LSTM, we need to define input weights and recurrent weights for the forget gates.

Implement LSTM

```python
def forward(self, x, h=None, c=None):
    """
    Inputs:
    - x: Input tensor (seq_len, batch_size, input_size)
    - h: Hidden vector (nr_layers, batch_size, hidden_size)
    - c: Cell state vector (nr_layers, batch_size, hidden_size)

    Outputs:
    - h_seq: Hidden vector along sequence (seq_len, batch_size, hidden_size)
    - h: Final hidden vetor of sequence(1, batch_size, hidden_size)
    - c: Final cell state vetor of sequence(1, batch_size, hidden_size)
    """
    # Below code handles the batches with varying sequence lengths
    lengths = None
    if isinstance(x, PackedSequence):
        x, lengths = pad_packed_sequence(x)

    # State initialization provided to you
    state_size = (1, x.size(1), self.hidden_size)
    if h is None:
        h = torch.zeros(state_size, device=x.device, dtype=x.dtype)
    if c is None:
        c = torch.zeros(state_size, device=x.device, dtype=x.dtype)
    assert state_size == h.shape == c.shape

    # Fill the following lists and convert them to tensors
    h_seq = []
    c_seq = []
```

```python
#############################################################
#  TODO: Perform the forward pass                           #
#############################################################

for xt in x.unbind(0):
    # Get the gates and update
    hiddens = self.W_hh(h) + self.W_xh(xt)
    gates, update = hiddens.split(
        (3 * self.hidden_size, self.hidden_size), dim=-1
    )
    gates = gates.sigmoid()
    update = update.tanh()

    # Update the hidden state
    fg, ig, og = gates.chunk(3, dim=-1)
    c = c * fg + update * ig
    h = og * c.tanh()

    # Keep the hidden state history
    h_seq.append(h)
    c_seq.append(c)

h_seq = torch.cat(h_seq, 0)
c_seq = torch.cat(c_seq, 0)

#############################################################
#                   END OF YOUR CODE                        #
#############################################################
```

We loop over the time step of the sequence and update parameters of the LSTM.

Questions? Piazza ☺